

1

Machine Learning

Ensemble learning

林嘉文 (Chia-Wen Lin)

清華大學電機系

cwlin@ee.nthu.edu.tw

No Free Lunch Theorem



No Free Lunch Theorem

- No Free Lunch Theorem: There is **NO** overall superior or inferior classification, if **no prior assumptions about the problem** are made
- Generate a group of base-learners which when combined has higher accuracy
- Different learners use different
 - Algorithms
 - Hyperparameters
 - Representations /Modalities/Views
 - Training sets
 - Subproblems
- Diversity vs accuracy

Bias & Variance

Unknown function f


Estimator $g_i = g(X_i)$ on sample X_i

Bias: $b_f(g) = \mathbb{E}[g] - f$

Variance: $\mathbb{E}[(g - \mathbb{E}[g])^2]$

Mean square error:

$$\begin{aligned} r(g, f) &= \mathbb{E}[(g - f)^2] \\ &= (\mathbb{E}[g] - f)^2 + \mathbb{E}[(g - \mathbb{E}[g])^2] \\ &= \text{Bias}^2 + \text{Variance} \end{aligned}$$

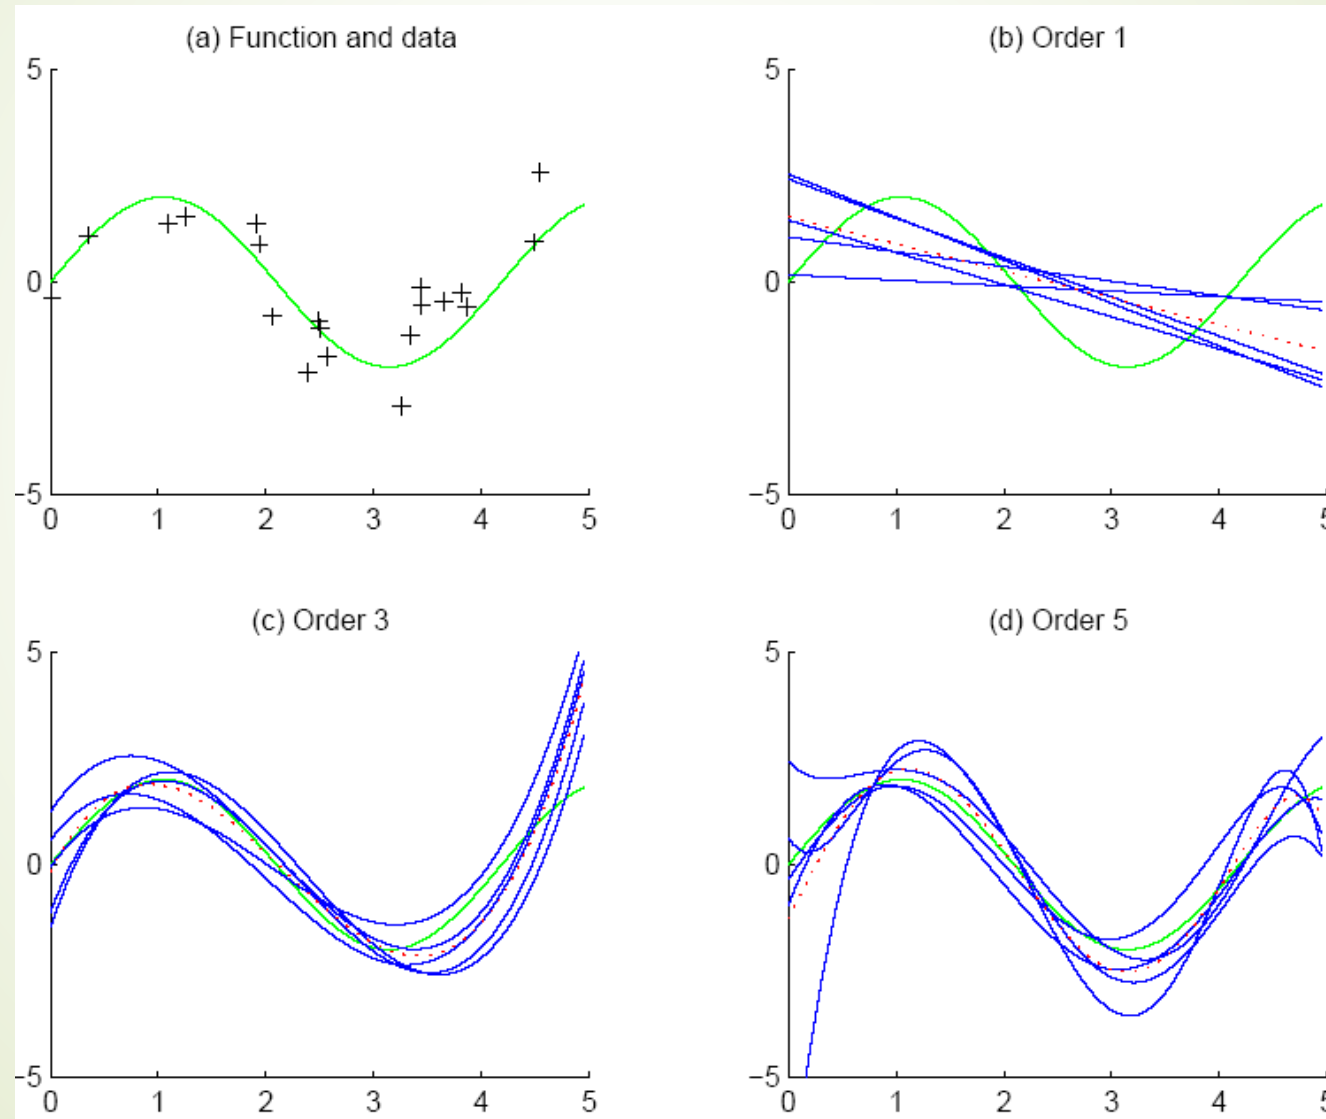


Bias/Variance Dilemma


$$\begin{aligned} & \mathbb{E}_D[(g(x, D) - f(x))^2] \\ &= \underbrace{(\mathbb{E}_D[g(x, D) - f(x)])^2}_{\text{bias}^2} + \underbrace{(\mathbb{E}_D[(g(x, D) - \mathbb{E}_D[g(x, D)])^2])}_{\text{variance}} \end{aligned}$$

- Bias: the difference between the expected value and the true value
- Variance: variations of estimated values
- Given training set D , as we increase complexity,
 - bias decreases (a better fit to data) and variance increases (fit varies more with data)
- High bias means usually low variance, and vice versa
- Bias/Variance dilemma: (Geman et al., 1992)


Bias/Variance Dilemma



Bias/Variance Dilemma



Ensemble Learning



Ensemble Learning: Intuition

- **Majority Vote**
- Suppose we have 5 completely independent classifiers each having 70% accuracy
 - $10 \cdot 0.7^3 \cdot 0.3^2 + 5 \cdot 0.7^4 \cdot 0.3^1 + 0.7^5 = 83.7\%$ majority accuracy
- 101 such classifiers
 - 99.9% majority vote accuracy

Ensemble Learning

- **Bagging** (Breiman 1994,...): Fit several classifiers to bootstrap **resampled** versions of the training data, and classify by majority vote.
- **Boosting** (Freund and Schapire 1995, Friedman et al. 1998,...): Fit many large or small classifiers to **reweighted** versions of the training data, then classify by weighted majority vote
- **Random forests** (Breiman 2001,...): Fancier version of bagging

Predict class label for unseen data by aggregating a set of predictions (classifiers learned from the training data)

Voting of Classifiers


- Linear combination

$$y = \sum_{j=1}^L w_j d_j$$

$$w_j \geq 0 \text{ and } \sum_{j=1}^L w_j = 1$$

- Classification

$$y_i = \sum_{j=1}^L w_j d_{ji}$$



Voting of Classifiers

Rule	Fusion function $f(\cdot)$
Sum	$y_i = \frac{1}{L} \sum_{j=1}^L d_{ji}$
Weighted sum	$y_i = \sum_j w_j d_{ji}, w_j \geq 0, \sum_j w_j = 1$
Median	$y_i = \text{median}_j d_{ji}$
Minimum	$y_i = \min_j d_{ji}$
Maximum	$y_i = \max_j d_{ji}$
Product	$y_i = \prod_j d_{ji}$

	C_1	C_2	C_3
d_1	0.2	0.5	0.3
d_2	0.0	0.6	0.4
d_3	0.4	0.4	0.2
Sum	0.2	0.5	0.3
Median	0.2	0.5	0.4
Minimum	0.0	0.4	0.2
Maximum	0.4	0.6	0.4
Product	0.0	0.12	0.032

Voting of Classifiers

- Bayesian perspective:

$$P(C_i|x) = \sum_{\text{all models } M_j} P(C_i|x, M_j) P(M_j)$$

If d_j are iid

$$E[y] = E\left[\sum_j \frac{1}{L} d_j\right] = \frac{1}{L} L \cdot E[d_j] = E[d]$$

$$\text{Var}(y) = \text{Var}\left(\sum_j \frac{1}{L} d_j\right) = \frac{1}{L^2} \text{Var}\left(\sum_j d_j\right) = \frac{1}{L^2} L \cdot \text{Var}(d_j) = \frac{1}{L} \text{Var}(d)$$

Bias does not change, variance decreases to $1/L$

- If dependent, error increases with positive correlation

$$\text{Var}(y) = \frac{1}{L^2} \text{Var}\left(\sum_j d_j\right) = \frac{1}{L^2} \left[\sum_j \text{Var}(d_j) + 2 \sum_j \sum_{i < j} \text{Cov}(d_i, d_j) \right]$$

Resampling for Classifier Design: Jackknife

- Remove some point from the training set: $D_{(i)}$
- Calculate the leave-one-out statistics with the new training set

$$\mu_{(i)} = \frac{1}{n-1} \sum_{j \neq i} x_j$$

- Repeat for all points
- Calculate the jackknife (leave-one-out) statistics


$$\mu_{(\cdot)} = \frac{1}{n} \sum_{i=1}^n \mu_{(i)}$$



Resampling for Classifier Design

- Arcing
 - Adaptive reweighting and combining
 - Reuse or select data in order to improve classification
 - **Bootstrap** data set
 - Randomly select n points from the training set D with replacement

Bagging



■ Bootstrap AGGregation

- Use multiple versions of a training set by drawing $n < N$ samples from D with replacement.
- Each of the bootstrap data sets is used to train a different component classifier
- The final classification decision is based on the votes of the component classifiers

Bagging

- Instability
 - A classifier/learning algorithm is called “unstable” if “small” changes in the training data lead to “large” changes in accuracy
 - In general, bagging improves recognition for unstable classifiers since it effectively averages over such discontinuities

Boosting

- Overview of boosting
 - Improve the accuracy of any given learning algorithm
 - First create a (**weak**) classifier with accuracy on the training set greater than average
 - Add new component classifiers to form an ensemble whose joint decision rule has arbitrarily high accuracy on the training set
 - The classification performance has been “boosted”

Boosting

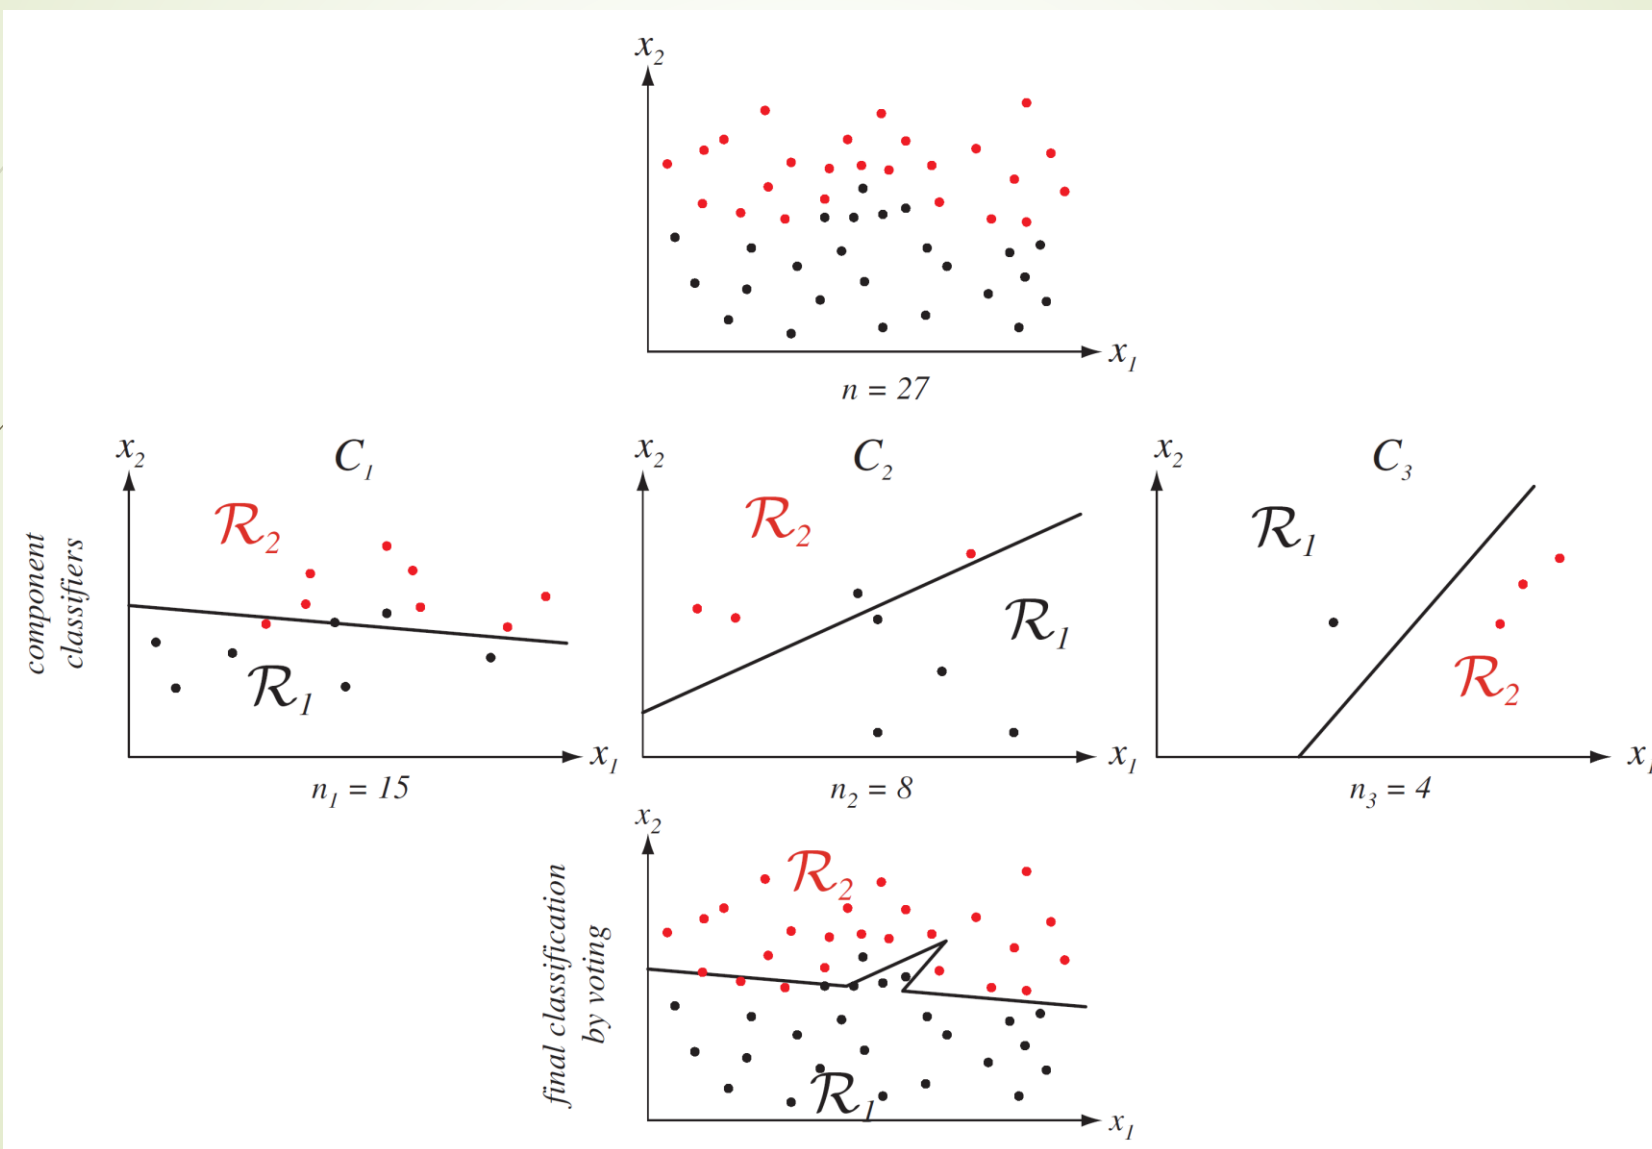
■ Boosting procedure

- Randomly select a set of $n_1 < n$ patterns from the full training set D **without replacement**; call this set D_1
- Train the first component classifier C_1 with D_1 ; C_1 is usually a weak learner
 - A weak learner has accuracy only slightly better than chance.
- Seek a second training set D_2 , that is “most informative” given component classifier C_1
 - Half of the patterns in D_2 should be correctly classified by C_1
 - **Half of the patterns in D_2 should be incorrectly classified by C_1 sampled from the remaining samples in D**
- Train the second component classifier C_2 with D_2

Boosting

- Boosting procedure
 - Then, the third component classifier C_3 with D_3
 - The samples in D_3 are selected from the remaining samples in D that are not well classified by voting by C_1 and C_2 (i.e., the classification results are not consistent)
 - The final classification decision is based on the votes of the component classifiers.

Boosting



Boosting

- How to decide n_1
 - We would like to train the ensemble classifier using all patterns in D .
 - A reasonable guess $n_1 \cong n_2 \cong n_3 = \frac{n}{3}$
 - In practice, we need to run the overall boosting procedure a few times, adjusting n_1 to use the full training set and get roughly equal partitions of the training set, if possible.

AdaBoosting

- Overview of AdaBoosting
 - A variant of boosting
 - Adaptive Boosting
 - Add weak learners until some desired low training error has been achieved.
 - Each training pattern receives a weight that determines its probability of being selected for a training set for an individual component classifier.
 - Adaboost “focuses on” the informative or “difficult” patterns.

AdaBoosting

Algorithm 1 (AdaBoost)

```

1 begin initialize  $\mathcal{D} = \{\mathbf{x}^1, y_1, \mathbf{x}^2, y_2, \dots, \mathbf{x}^n, y_n\}, k_{max}, W_1(i) = 1/n, i = 1, \dots, n$ 
2    $k \leftarrow 0$ 
3   do  $k \leftarrow k + 1$ 
4     Train weak learner  $C_k$  using  $\mathcal{D}$  sampled according to distribution  $W_k(i)$ 
5      $E_k \leftarrow$  Training error of  $C_k$  measured on  $\mathcal{D}$  using  $W_k(i)$ 
6      $\alpha_k \leftarrow \frac{1}{2} \ln[(1 - E_k)/E_k]$ 
7      $W_{k+1}(i) \leftarrow \frac{W_k(i)}{Z_k} \times \begin{cases} e^{-\alpha_k} & \text{if } h_k(\mathbf{x}^i) = y_i \text{ (correctly classified)} \\ e^{\alpha_k} & \text{if } h_k(\mathbf{x}^i) \neq y_i \text{ (incorrectly classified)} \end{cases}$ 
8   until  $k = k_{max}$ 
9   return  $C_k$  and  $\alpha_k$  for  $k = 1$  to  $k_{max}$  (ensemble of classifiers with weights)
10 end


```

$$g(\mathbf{x}) = \sum_{k=1}^{k_{\max}} \alpha_k h_k(\mathbf{x}),$$


$$h_k(\mathbf{x}) = +1 \text{ or } -1,$$

the category label given to \mathbf{x} by component classifier C_k

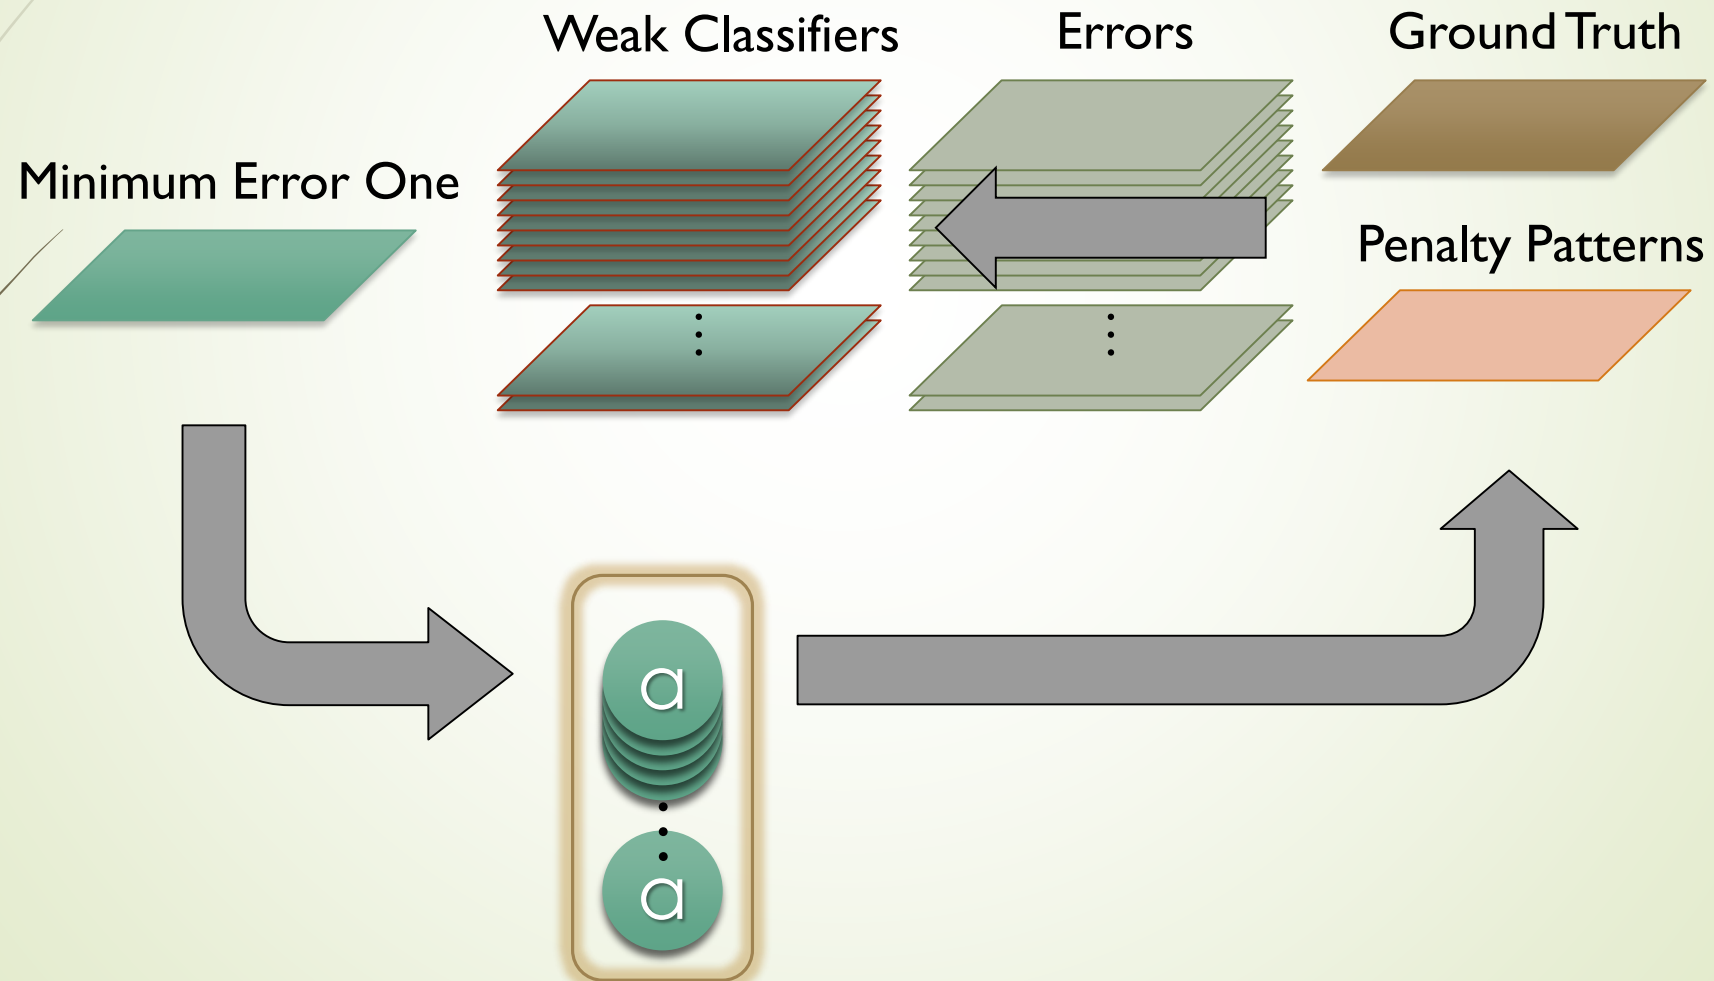
AdaBoosting




AdaBoosting for Classification



AdaBoosting for Classification



AdaBoosting for Classification




Random Forest

- Decision Tree Bagging
 - Given a training set $X = x_1, \dots, x_n$ with labels $Y = y_1, \dots, y_n$, bagging repeatedly (L times) selects a random sample with replacement of the training set and fits trees to these samples
 - For $i = 1, \dots, L$:
 - Sample, with replacement, m training examples from X, Y ; call these X_i, Y_i
 - Train a decision or regression tree f_i on X_i, Y_i
 - Prediction for an unseen x' can be made by averaging the predictions from all the individual trees on x'
- Random Forest use a random subset of the features in tree learning, called "feature bagging"

Cross-Validation

- Simple validation
 - Split the set of labeled training samples D into two parts
 - Traditional training set
 - Train the classifier
 - Validation set
 - Estimate the generalization error

Cross-Validation



Cross-Validation

- m -fold cross-validation
 - Split the training set D into m disjoint sets of equal size n/m
 - The classifier is trained m times, each time with a different set held out as a validation set
 - The estimated performance is the mean of the m errors
 - Leave-one-out approach ($m = n$)

Cross-Validation

- 5-fold cross-validation -> split the training data into 5 equal folds
- 4 of them for training and 1 for validation

